

# Современные нейросетевые технологии

---

Лекция 6. Процесс обучения  
сетей

1. Препроцессинг данных и инициализация весов
2. Поиск гиперпараметров
3. Продвинутая регуляризация
4. Transfer learning

[github.com/balezz/modern\\_dl](https://github.com/balezz/modern_dl)

Срок сдачи А6 – 08.10.2022 г.

Источники:

- [dlcourse.ai](https://dlcourse.ai)
- [cs231n.stanford.edu](https://cs231n.stanford.edu)
- [cs230.stanford.edu](https://cs230.stanford.edu)


## Neural Networks

Linear score function:

$$f = Wx$$


2-layer Neural Network

$$f = W_2 \max(0, W_1 x)$$



$$\frac{dL}{dx} = \frac{dh}{dx} \cdot \frac{ds}{dh} \cdot \frac{dL}{ds}$$

# Learning network parameters through optimization




```
# Vanilla Gradient Descent

while True:
    weights_grad = evaluate_gradient(loss_fun, data, weights)
    weights += - step_size * weights_grad # perform parameter update
```

# Activation Functions


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



## Convolutional Layer




1. Определите препроцессинг данных.
2. Выберите простую архитектуру сети.
3. Проверьте loss на необученных весах. (e.g. 2.3 for 10 classes)
4. Проверьте, что сеть переобучается на малой выборке (20 samples per class).

```
Finished epoch 195 / 200: cost 0.002694, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 196 / 200: cost 0.002674, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 197 / 200: cost 0.002655, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 198 / 200: cost 0.002635, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 199 / 200: cost 0.002617, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 200 / 200: cost 0.002597, train: 1.000000, val 1.000000, lr 1.000000e-03
finished optimization. best validation accuracy: 1.000000
```

# Препроцессинг данных

`tf.keras.layers.Normalization()`





## Batch Normalization

[Ioffe and Szegedy, 2015]

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Note: at test time BatchNorm layer functions differently:**

The mean/std are not computed based on the batch. Instead, a single fixed empirical mean of activations during training is used.

(e.g. can be estimated during training with running averages)

- First idea: **Small random numbers**  
(gaussian with zero mean and  $1e-2$  standard deviation)




```
W = 0.01* np.random.randn(D, H)
```

Works ~okay for small networks, but problems with deeper networks.


# Инициализация весов

```
input layer had mean 0.000927 and std 0.998388
hidden layer 1 had mean -0.000117 and std 0.213081
hidden layer 2 had mean -0.000001 and std 0.047551
hidden layer 3 had mean -0.000002 and std 0.018630
hidden layer 4 had mean 0.000001 and std 0.002378
hidden layer 5 had mean 0.000002 and std 0.000532
hidden layer 6 had mean -0.000000 and std 0.000119
hidden layer 7 had mean 0.000000 and std 0.000026
hidden layer 8 had mean -0.000000 and std 0.000006
hidden layer 9 had mean 0.000000 and std 0.000001
hidden layer 10 had mean -0.000000 and std 0.000000
```

Forward pass:  
layers outputs becomes to zeros  
Backward pass:  
gradients are zeros too




# Инициализация весов




## Initialization too small:

Activations go to zero, gradients also zero.  
No learning



## Initialization too big:

Activations saturate (for tanh),  
Gradients zero, no learning



## Initialization just right:

Nice distribution of activations at all layers.  
Learning proceeds nicely

# Инициализация весов

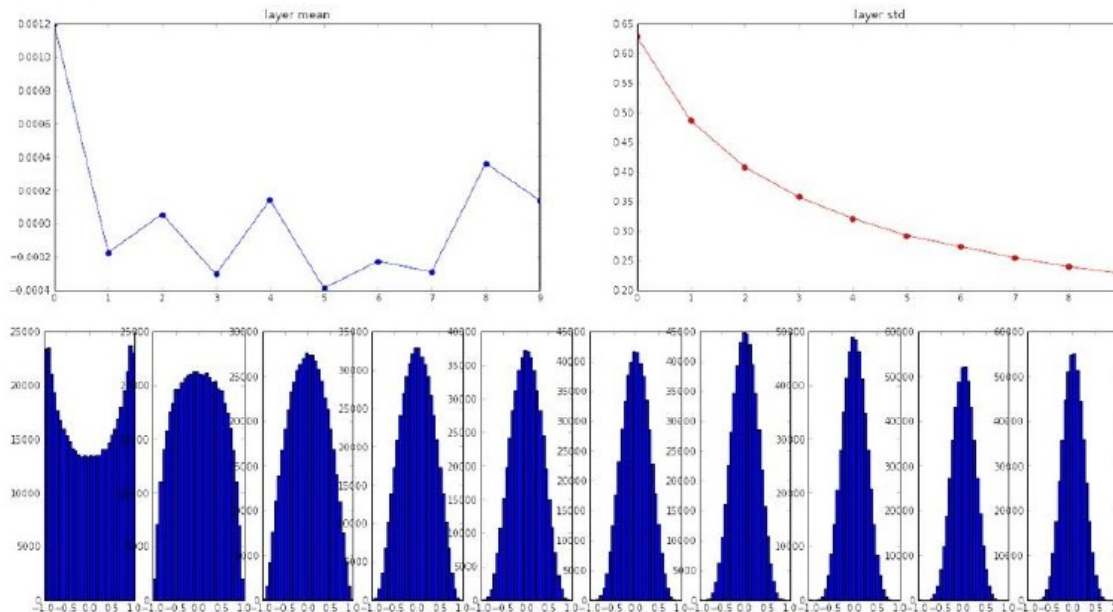
This works with tanh activations

```
input layer had mean 0.001800 and std 1.001311
hidden layer 1 had mean 0.001198 and std 0.627953
hidden layer 2 had mean -0.000175 and std 0.486051
hidden layer 3 had mean 0.000055 and std 0.407723
hidden layer 4 had mean -0.000306 and std 0.357108
hidden layer 5 had mean 0.000142 and std 0.320917
hidden layer 6 had mean -0.000389 and std 0.292116
hidden layer 7 had mean -0.000228 and std 0.273387
hidden layer 8 had mean -0.000291 and std 0.254935
hidden layer 9 had mean 0.000361 and std 0.239266
hidden layer 10 had mean 0.000139 and std 0.228008
```

```
W = np.random.randn(fan_in, fan_out) / np.sqrt(fan_in) # layer initialization
```

"Xavier initialization"  
[Glorot et al., 2010]

**Reasonable initialization.**  
(Mathematical derivation  
assumes linear activations)






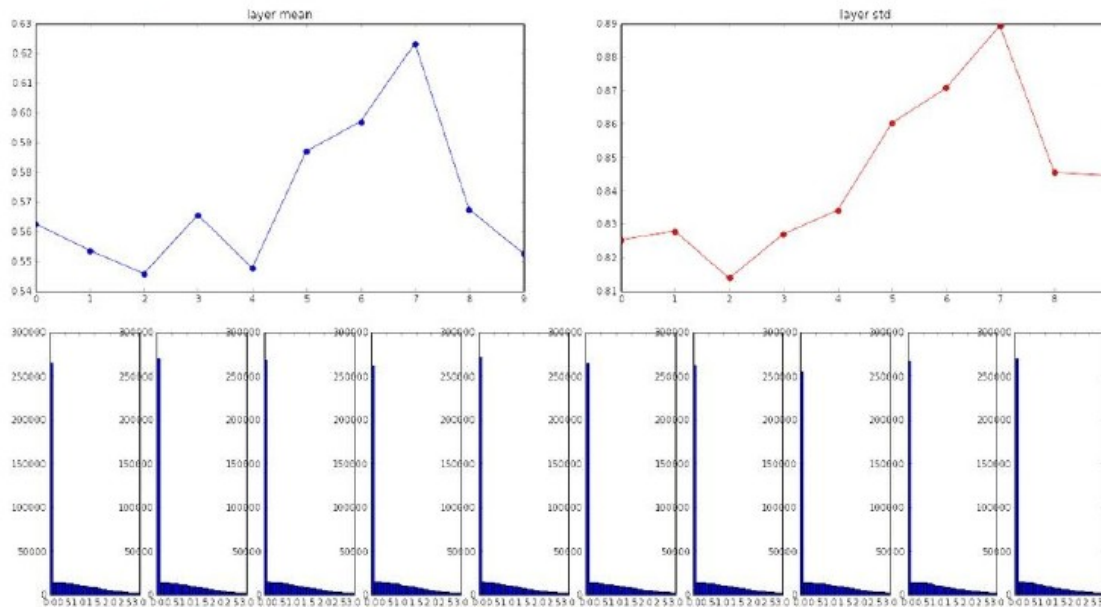
# Инициализация весов


This works with ReLU activations

input layer had mean 0.000581 and std 0.999444  
hidden layer 1 had mean 0.562488 and std 0.825232  
hidden layer 2 had mean 0.553614 and std 0.827835  
hidden layer 3 had mean 0.545867 and std 0.813855  
hidden layer 4 had mean 0.565396 and std 0.826902  
hidden layer 5 had mean 0.547678 and std 0.834092  
hidden layer 6 had mean 0.587103 and std 0.860035  
hidden layer 7 had mean 0.596867 and std 0.870610  
hidden layer 8 had mean 0.623214 and std 0.889348  
hidden layer 9 had mean 0.567498 and std 0.845357  
hidden layer 10 had mean 0.552531 and std 0.844523

```
W = np.random.randn(fan_in, fan_out) / np.sqrt(fan_in/2) # layer initialization
```

He et al., 2015  
(note additional /2)





## Random Search vs. Grid Search

*Random Search for  
Hyper-Parameter Optimization  
Bergstra and Bengio, 2012*

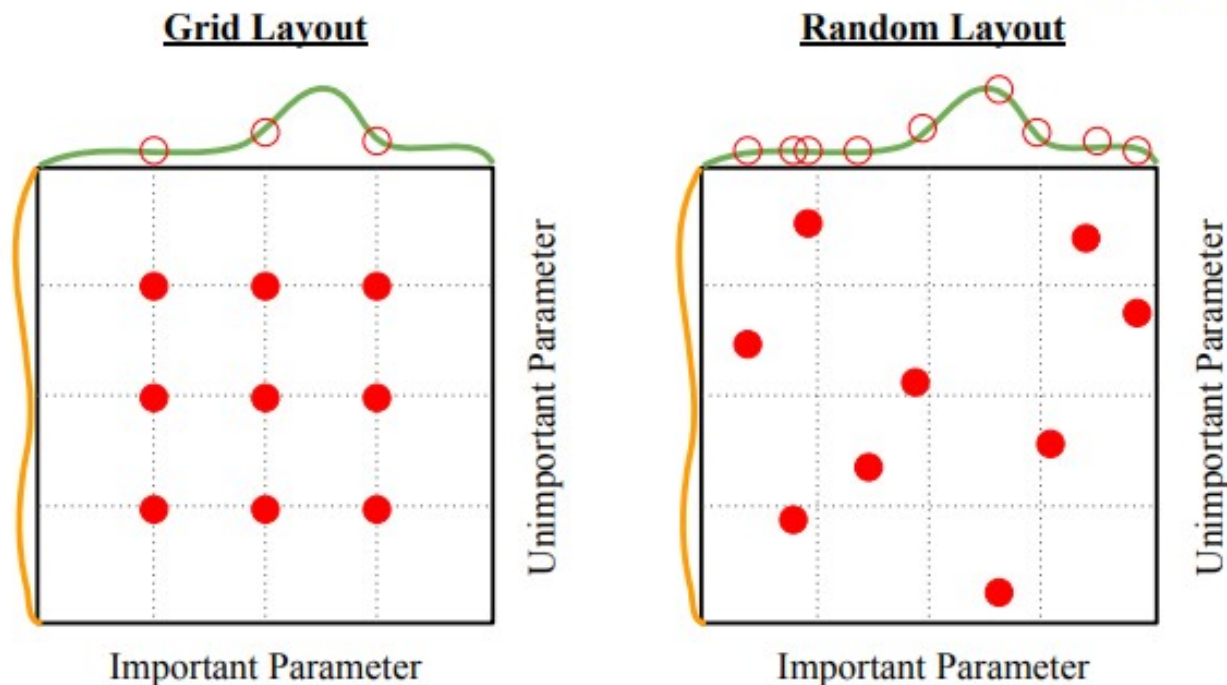




Illustration of Bergstra et al., 2012 by Shayne  
Longpre, copyright CS231n 2017



# Beyond Training Error





Better optimization algorithms  
help reduce training loss



But we really care about error on new  
data - how to reduce the gap?

## Early Stopping





Stop training the model when accuracy on the validation set decreases

Or train for a long time, but always keep track of the model snapshot that worked best on val


# Regularization: Dropout

In each forward pass, randomly set some neurons to zero  
Probability of dropping is a hyperparameter; 0.5 is common



# Regularization: Dropout

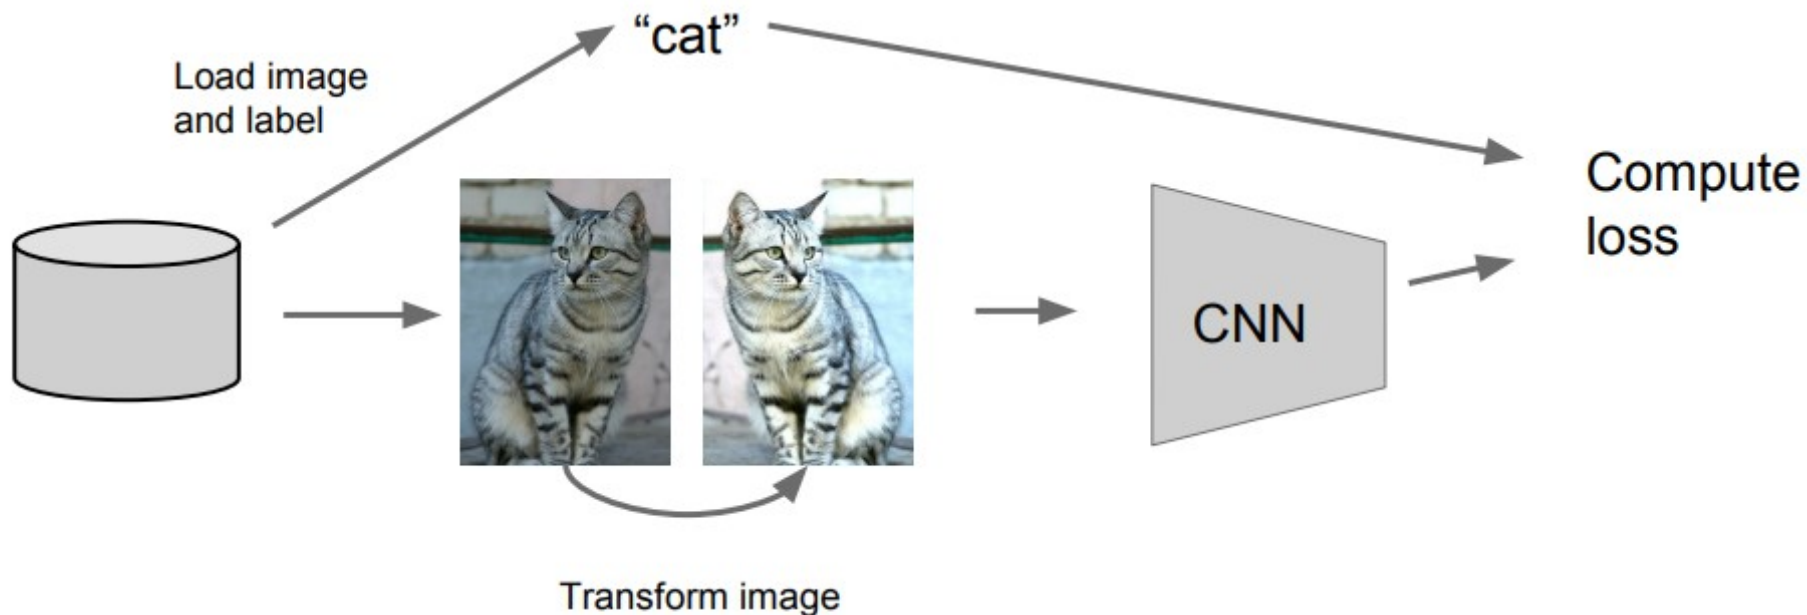
How can this possibly be a good idea?



Forces the network to have a redundant representation;  
Prevents co-adaptation of features



# Regularization: Data Augmentation





# Data Augmentation

## Horizontal Flip



## Color Jitter

Simple: Randomize  
contrast and brightness



# Data Augmentation

## Random crops and scales

**Training:** sample random crops / scales


ResNet:

1. Pick random  $L$  in range  $[256, 480]$
2. Resize training image, short side =  $L$
3. Sample random  $224 \times 224$  patch

**Testing:** average a fixed set of crops


ResNet:

1. Resize image at 5 scales:  $\{224, 256, 384, 480, 640\}$
2. For each size, use 10  $224 \times 224$  crops: 4 corners + center, + flips

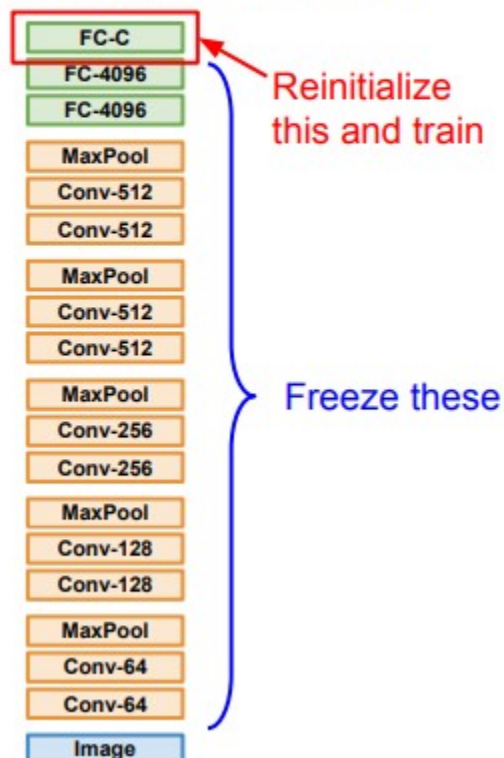


## Transfer Learning with CNNs

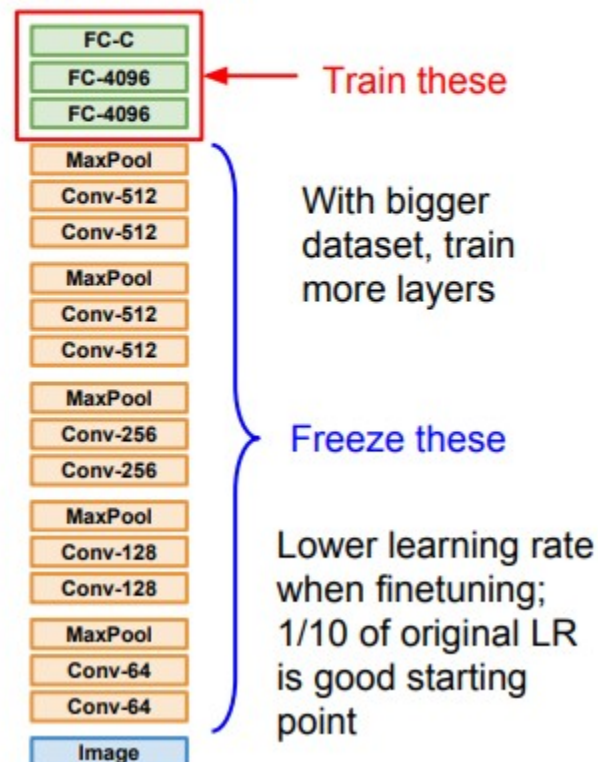
### 1. Train on Imagenet



### 2. Small Dataset (C classes)



### 3. Bigger dataset



Donahue et al, "DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition", ICML 2014  
Razavian et al, "CNN Features Off-the-Shelf: An Astounding Baseline for Recognition", CVPR Workshops 2014